



14-848 Cloud Infrastructure: Design, Analysis and Implementation

Meeting Dates/Times: Tue/Thu 3:30 - 4:50 PM ET (12:30 – 1:50 PM PT)

Location: CIC 1201 for PIT Students & B23 109 for SV Students

Course Webpage: <https://canvas.cmu.edu/courses/45143>

Semester: Spring, **Year:** 2025

Units: 12, **Sections:** A (Pittsburgh), SV (Silicon Valley)

Instructor

Name	Mohamed Farag (he/his/him) – Addressed as “Mohamed” or “Dr. Farag”
Contact Info	farag@cmu.edu
Office Hours Location	CMU Remote
Office Hours	Monday/Wednesday 11:30am – 12:30PM ET (8:30 – 9:30am PT) Conducted remotely via Zoom. URL: https://cmu.zoom.us/j/94117627561?pwd=mmJpBtil76BwNQUDZ4KUdWoFbRpvOa.1

Course Description

- Cloud computing has emerged as a new paradigm for efficient and highly elastic delivery of computing services over the Internet to achieve economies of scale. It focuses on the delivery of services via on-demand and fluidly scalable shared resources. This course focuses on the design and implementation of networked systems and software necessary to implement the infrastructure for elastic, global-scale computing, and storage clouds. Upon successful completion of the course, students will be able to design, implement, and analyze the infrastructure underlying cloud and edge-based services. Topics covered in this course include data center networking, cloud-scale storage, virtualization, computing programming models, large-scale distributed computing, edge computing, resource utilization and sharing, and cloud service frameworks. The course material will focus on recent and landmark research papers and existing tools and software systems. Students will have substantial programming project work in which they design, implement, and analyze aspects of cloud infrastructure and services. Students are expected to be proficient in object-oriented programming and Linux system programming and command-line tools.
- The premise of this course is to build a broad and solid foundation in cloud computing that will pay significant dividends throughout a student’s research and work career across cloud computing and DevOps related fields. In this class, we will focus on five main themes:
 - Cloud computing concepts and programming paradigms
 - Cloud data storage
 - Cloud Deployment and Orchestration
 - Big data processing frameworks.
 - Miscellaneous topics related to cloud computing.

- **Prerequisites:** project course in software systems (e.g., networking, OS, etc.) and object-oriented programming experience should be sufficient
- **Class Structure:** Lectures will be synchronous, and attendance is required. Lecture slides are delivered via [TopHat](#), an online course delivery system. TopHat can be accessed on a smartphone or laptop. For this purpose, I will ask you in advance to bring your laptop or mobile device, and we will review how to use the tool together in class. If you do not have the necessary equipment, please contact your [HUB liaison](#) who is available to help you tap into appropriate resources. Join code is posted on the course Canvas page. Recordings will be made available to students for asynchronous access. Class meeting times will include some lectures, quizzes, discussion of course materials and ongoing project tasks, and other activities as appropriate. Class sessions may be partitioned/rotated as needed to accommodate reduced classroom capacity due to COVID-19 restrictions.
- This class doesn't restrict students to use specific programming language but the lecture materials target Python programming language. In general, comparable object-oriented programming experience is required and you will be asked to do several tasks in your favorite programming language to gain the valuable perspective that this class offers. If you have questions about your programming experience, please let me know at the beginning of the semester, so we can decide how to proceed.

Learning Objectives

In taking and successfully completing this course, students will:

- Understand the current cloud computing technology and critical technology trends that are enabling cloud computing, including cloud Computing architecture along with the services and the applications that cloud computing offers;
- Gain hands-on experience in building and analyzing cloud infrastructure components and developing and testing services that operate within them;
- Develop big-data processing applications to handle large loads of data; and
- Simplify the deployment and portability of large complex applications to promote continuous integration and continuous delivery concepts using Docker.

The work that you will carry out in this course maps to some of the learning outcomes that INI has defined as vital to the full CMU graduate-level experience. Specifically:

- self-directed learning
- critical thinking; and
- complex problem solving.

Learning Resources

There is no primary textbook, as most reading material will come from research papers and other technical documentation. Additional background reading material can be suggested upon request.

Important Dates

- January 20th, 2025: MLK Jr. Day (no office hours)
- March 3rd – March 7th, 2025: Spring break (no classes/office hours).
- April 3rd – April 5th 2025: Spring Carnival (no classes/office hours).
- April 24th, 2025: Final exam (during lecture time)

Assessments

Students are encouraged to attend class regularly, read the assigned reading material and participate in class discussions. The final grade will be based upon 1 exam, 1 project, 7 homework assignments, and in-class quizzes.

Final Exam	Project	Assignments	Quizzes
15%	20%	40%	25%

- **Quizzes:** are offered during each lecture via Canvas. Each quiz will be accessible via a unique access code that will be provided to the students. Students will have 2-3 minutes to answer 1-2 multiple choice questions.
 - Please note that students are not allowed to share Quiz Access Code with their peers who didn't attend the class in-person without pre-approval from the instructor. Sharing the quiz access code, without instructor's approval, is considered an academic integrity violation.
- **Assignments:** will provide the opportunity to practice the concepts that are taught during the lectures. Students are expected to spend a good amount of time on their own to learn implementation details that are not provided during the lectures. Students will receive 7 assignments throughout the semester. Assignment schedule is shown on the last section of this syllabus. **While we will use Google Cloud Platform for most assignments, students will have to apply for trial version to use Azure and AWS. The trial version will require the students to enter their credit card information without being billed.**
- **Project:** details are released in week 3. Each student will have the option to choose another student for the project and you will choose one of two project options to submit. Students will be expected to record a video including a code-walkthrough of their work and functionality demo showing the running version of their application. Project submission deadline is **April 2nd, 2025 11:59PM ET /8:59pm PT.** Course project will include a lot of self-learning that is needed from the students to complete the project. The course instructor will provide project-related hints, high-level directions, and clarifications during the lectures. However, students shouldn't expect any additional project support during office hours or via emails. Students are highly encouraged to give themselves enough time to learn the skills they need to complete the project. There will be a checkpoint to ensure that students are making good continuous progress (refer to the proposed course schedule in the last section of this syllabus). Project grading rubric and evaluation will be released along with project details.
- **Final Exam:** is an open-note test.
 - Students will have access to all the **PDFs for lectures, readings and HW solutions.** Students can **bring any hard-copied materials with them.**
 - Students are required to follow the schedule of their registered section. **On the scheduled final lecture of each section, final exam will be released only to the registered students of the corresponding section.** Each section will have its final exam version(s).
 - Exam will be offered via **Lockdown Browser** and **no knowledge exchange is allowed among students during the exam.**
 - Students are expected to install and test Lockdown browser on their machines ahead of the exam. If students face an issue with Lockdown browser installation, students must reach out to the instructors **no later than 2 weeks** before the final exam date.
 - **Sharing hard-copied notes is prohibited during the exam.**
 - **You can get full score on the final exam if you obtain TWO of the following certifications two weeks before the final exam:**

- [Google Cloud: Associate Cloud Engineer](#)
 - [AWS Certified Solutions Architect – Associate](#)
 - [Microsoft Certified: Azure Developer Associate](#)
 - [GCP Professional Data Engineer](#)
 - [AWS Certified Data Engineer - Associate](#)
 - [Microsoft Azure Data Engineer Associate](#)
 - [Certified Kubernetes Application Developer](#)
- Late acquisition of these certificates beyond the deadline won't be accepted. Students who would like to waive their final exam must upload a proof of obtaining two certificates two weeks prior to the exam. Late submissions WILL NOT be accepted.
 - Obtaining only 1 certificate of the 3 won't provide any partial benefit on the final exam score

Students will be assigned the following final letter grades, based on rubric provided in the above table. +/- are assigned in equal intervals to provide further granularity.

Grade	Percentage Interval
A/A-	[85-100%], A starts from 93
B	[70-85%)
C	[55-70%)
D	[40-55%)
R (F)	Below 40%

Homework, Course Project, Grace Days, Late Submission & Grading Policies

- Students can submit the assignments online submitting their **GitHub Org URL via Gradescope**. For this purpose, sign up for a (free) account at GitHub.
- **Students must complete their GitHub code commits and the submission of their GitHub URL over Gradescope before the submission deadline.**
- Each assignment has **1 grace day beyond the assignment deadline**.
- The course project has **1 grace day beyond the project deadline**.
- Grade days **are not transferable**.
- Students are strongly encouraged to complete the assignments as early as possible. No extensions for homework will be offered beyond the grace day, even in the case of technical difficulties.
- **Students will have 3 days to submit an assignment after the grace day extension** and a late penalty will be applied. Late penalties are applied based on the timestamp of the last code commit on GitHub and it will follow this equation:
 - 5 points for up to 24 hours delay
 - 15 points for the next 24 hours delay
 - 25 points for the next 24 hours delay
 - 100 points penalty (no grade) after this time.
- **Late submissions for the course project will receive no grade (0 points).**

- After homework and project grades are released, **Regrade requests can be made for 24 hours via Gradescope**. Regrade requests **cannot be submitted via email**. Also, regrade requests **cannot be submitted after the 24 hours window is closed**.
- For grading inquiries, email the TA and copy the instructor on your email.
- Students are encouraged to read their privacy rights in [Family Educational Rights and Privacy Act \(FERPA\)](#).

Guidelines for Office Hours and Out-of-class Questions

- If you have materials-related questions, you can post them on **Piazza** or ask them during office hours. You should plan to receive the question responses on **Piazza** during office hours' allocated time.
- You may email personal inquiries and severe emergencies to the instructor(s). The email subject line should begin with "**14-848**". Emails sent to the instructor should be 2-3 lines maximum and the instructor's response will not be more than a few words. Generally, the course instructor is available to respond to personal inquiries during M-F 9am-8pm ET. Please don't expect responses over the weekend or late at night.

Expectations for Class Attendance and Class Absence

- Classes are offered **synchronously and in-person**.
- Class attendance and participation are important components of the learning in this course. To account for this, a portion of the final grade is based on quizzes that are offered during the lectures (see assessment section). That said, I also recognize that students may need to miss class for a variety of reasons (religious observance, job interview, university-sanctioned event, or illness). **For that reason, all students are permitted two class absences (including their quizzes) without any impact on the final grade.** If you encounter extenuating circumstances and must miss more than two classes, please email the course instructor and follow the student affairs guidance below:
 "Based on the Student Affairs recommendation for guiding students who may be absent for more than a few classes due to a personal issue (injury/illness or other), the student affairs recommend that students either contact the office of disability resources (particularly if it may be an extended issue) or fill out a CMU cares referral form for the student to get them in touch with a student affairs liaison. The student can also be advised to fill out the form on their own behalf if they are seeking guidance on an issue they are dealing with. Here is the form: <https://www.cmu.edu/wellbeing/resources/student-support-resources.html>"
- When attending the class in-person, I expect that you will abide by all behaviors indicated in [A Tartan's Responsibility](#), including any timely updates based on the current conditions.

Recording of Class Sessions

- Class recordings will be available after each lecture. Please note it may take a few hours for the recording to become available. The Class recording catalog will be published on Canvas.

Academic Integrity

- Discussing assignments with your classmates is allowed and encouraged, but it is important that every student gets practice working on these problems. This means that **all** the work you turn in must be your own. You must devise and write your own solutions and carry out your own tests. The general policy on homework collaboration is:
 - You may use material that we explicitly provide you for the assignment. No attribution is required.
 - You may use other course material, including lectures, Piazza posts by the instructors, and material from Canvas. For any such use involving code, you must provide clear attribution, indicating the source, and where the included material begins and ends.
 - Beyond this, you may not obtain code or other solution information, either by copying, retyping, or looking at files or documents from any of the above-listed courses for this semester, or a previous

semester.

- You may not obtain code or other solution information from an unauthorized external source, including web pages, code repositories, blog posts, etc
 - You may not look at someone else's code (or other documents.) This includes one person looking at code and describing it to another. There is no notion of looking "too much," since no looking is allowed at all.
 - You may not make use of any information about the assignments posted online, except for the authorized sources provided with the assignment document.
 - You may get assistance on an assignment from the instructors, course staff, and university tutors.
 - You may only get high-level, strategic advice from others, including current and former students, and people external to the university. Forbidden forms of advice include: anything more detailed than a brief verbal description or block diagram, any kind of code or pseudo-code, explicit directions on how to assemble allowed blocks of code, and code-level debugging assistance.
 - You may not provide detailed help with an assignment to students this semester or in future semesters for any of the above-listed courses (unless you are serving as a teaching assistant or instructor for the course).
 - You may not supply a copy of a file or document to an individual student or via a public channel, such as a blog post.
 - You may not have any of your solution files in unprotected directories or in unprotected code repositories, either by putting files in an unprotected location or by allowing protections to lapse. Be sure to store your work in protected directories, and log off when you leave an open cluster, to prevent others from copying your work. If you make use of a code repository, such as GitHub, make sure your work is kept private, even after you have left CMU.
 - You may not provide electronic, verbal, or written descriptions of code or other solution information.
 - You may clarify ambiguities or vague points in class handouts or textbooks.
- Students may use Generative AI platforms (e.g., ChatGPT) to assist them with understanding the homework concepts. However, students can't use the direct homework question in their prompt to the AI tool. In addition, students are expected to cite the text (or code) that was generated from ChatGPT carefully. This includes scenarios where generative-AI was used to generate base-code/scenario and minor (or significant) changes have been made to it.
 - Students may not receive assistance from Generative AI platforms that would count for more than 30% of the homework solution.
 - Please talk to the instructor if you have any questions about this policy. Any form of plagiarism or cheating will result in sanctions to be determined by the instructors, including grade penalties (such as negative points for the assignment or reductions in letter grade) or course failure. Students taking the course pass/fail may have this status revoked. We are also obliged in report violations to your academic program and the appropriate University authorities. Please refer to the [University Policy on Academic Integrity](#).

Student Wellness

- The last few years have been challenging. We are all under a lot of stress and uncertainty at this time. I encourage you to find ways to move regularly, eat well, and reach out to your support system or me (farag@cmu.edu) if you need to. We can all benefit from support in times of stress, and this semester is no exception.

Diversity Statement

- **We must treat every individual with respect.** We are diverse in many ways, and this diversity is fundamental to building and maintaining an equitable and inclusive campus community. Diversity can refer to multiple ways that we identify ourselves, including but not limited to race, color, national origin, language, sex, disability, age,

sexual orientation, gender identity, religion, creed, ancestry, belief, veteran status, or genetic information. Each of these diverse identities, along with many others not mentioned here, shape the perspectives our students, faculty, and staff bring to our campus. We, at CMU, will work to promote diversity, equity and inclusion not only because diversity fuels excellence and innovation, but because we want to pursue justice. We acknowledge our imperfections while we also fully commit to the work, inside and outside of our classrooms, of building and sustaining a campus community that increasingly embraces these core values.

- Each of us is responsible for creating a safer, more inclusive environment.
- Unfortunately, incidents of bias or discrimination do occur, whether intentional or unintentional. They contribute to creating an unwelcoming environment for individuals and groups at the university. Therefore, the university encourages anyone who experiences or observes unfair or hostile treatment on the basis of identity to speak out for justice and support, within the moment of the incident or after the incident has passed. Anyone can share these experiences using the following resources:
 - **Center for Student Diversity and Inclusion:** csdi@andrew.cmu.edu, (412) 268-2150
 - **Report-It online anonymous reporting platform:** reportit.net username: tartans password: plaid
- All reports will be documented and deliberated to determine if there should be any following actions. Regardless of incident type, the university will use all shared experiences to transform our campus climate to be more equitable and just.

Food Insecurity

- If you are worried about affording food or feeling insecure about food, there are resources on campus that can help. Any undergraduate or graduate student can visit the CMU Pantry and receive food for free. Follow the directions on the [CMU Pantry website](#) to schedule your visit.

Disability Resources

- If you have a disability and have an accommodations letter from the [Disability Resources office](#), we encourage you to discuss your accommodations and needs with us as early in the semester as possible. We will work with you to ensure that accommodations are provided as appropriate. If you suspect that you may have a disability and would benefit from accommodations but are not yet registered with the Office of Disability Resources, we encourage you to contact them at access@andrew.cmu.edu.

Religious Observances

- In order to accommodate the observance of religious holidays, students should inform the instructor by email, within the first two weeks of the term, of any such days which conflict with scheduled class activities.

Student Academic Success Center

SASC programs to support student learning include the following (program titles link to webpages):

- [Academic Coaching](#): This program provides holistic, one-on-one peer support and group workshops to help undergraduate and graduate students implement habits for success. Academic Coaching assists students with time management, productive learning and study habits, organization, stress management, and other skills. Request an initial consultation here.
- [Peer Tutoring](#): Peer Tutoring is offered in two formats for students seeking support related to their coursework. Drop-In tutoring targets our highest demand courses through regularly scheduled open tutoring sessions during the fall and spring semesters. Tutoring by appointment consists of ongoing individualized and small group sessions. You can utilize tutoring to discuss course related content, clarify and ask questions, and work through practice problems.
- [Communication Support](#): Communication Support offers free one-on-one communication consulting as well

as group workshops to support strong written, oral, and visual communication in texts including IMRaD and thesis-driven essays, data-driven reports, oral presentations, posters and visual design, advanced research, application materials, grant proposals, business and public policy documents, data visualisation, and team projects. Appointments are available to undergraduate and graduate students from any discipline at CMU. Schedule an appointment (in-person or video), attend a workshop, or consult handouts or videos to strengthen communication skills.

- [Language and Cross-Cultural Support](#): This program supports students seeking help with language and cross-cultural skills for academic and professional success through individual and group sessions. Students can get assistance with writing academic emails, learning expectations and strategies for clear academic writing, pronunciation, grammar, fluency, and more. Make an appointment with a Language Development Specialist to get individualized coaching.
- [Supplemental Instruction \(SI\)](#): This program offers a non-remedial approach to learning in historically difficult courses at CMU. It utilizes a peer-led group study approach to help students succeed and is facilitated by an SI leader, a CMU student who has successfully completed the course. SI offers a way to connect with other students studying the same course, a guaranteed weekly study time that reinforces learning and retention of information, as well as a place to learn and integrate study tools and exam techniques specific to a course.

Mental Health Resources

- [CaPS has partnered with TimelyCare](#) for provision of virtual well-being services, including immediate emotional support 24/7 as frequently as needed, scheduled appointments with therapists that can be chosen by identity group and other features, health coaching (e.g., sleep issues, weight management, etc.), and group sessions for things like yoga, meditation, etc.

Preliminary Course Schedule (Subject to Change)

Date	Topic	Notes
Week-1 (Jan. 13 th)	- Introduction & Syllabus - Virtualization Basics	- System Setup homework released
Week-2 (Jan. 20 th)	- Containerization - Lab: Containerization	- System Setup homework deadline. - Docker homework released
Week-3 (Jan. 27 th)	- Deployment Orchestration - Pokémon Go Case Study - Lab: Deployment Orchestration	- Docker homework deadline - Course Project released - Kubernetes homework is released
Week-4 (Feb. 3 rd)	- Kafka - Lab: Confluent Kafka	
Week-5 (Feb. 10 th)	- Infrastructure-as-a-Code - Infrastructure-as-a-Code (Cont'd) Terraform	- Kubernetes homework deadline - Kafka Homework released
Week-6 (Feb. 17 th)	- Lab: Terraform - Cloud Infrastructure Concepts	
Week-7 (Feb. 24 th)	- Cloud Data Storage Models - Lab: NoSQL Database - Neo4j AuraDB	- Course Project Checkpoint
Spring Break (Mar. 3 rd - Mar. 7 th)		
Week-8 (Mar. 10 th)	- Introduction to Hadoop - Hadoop HDFS	- Kafka Homework deadline - Terraform Homework released

Week-9 (Mar. 17 th)	- Hadoop MapReduce	- Terraform Homework deadline - Hadoop MapReduce homework is released.
Week-10 (Mar. 24 th)	- Big Data Algorithms - Introduction to Spark	- Hadoop MapReduce homework deadline
Week-11 (Mar. 31 st)	- Spark (Cont'd)	- Course project submission deadline
Week-12 (Apr. 7 th)	- Lab: Spark Programming - Metaverse, Edge Computing and Fog Computing	- Apache Spark homework released
Week-13 (Apr. 14 th)	- Cloud Security & Privacy Concepts - DevSecOps	- Apache Spark homework deadline
Week-14 (Apr. 21 st)	- Introduction to GenAI on the Cloud using Vertex AI - Final Exam	